

LLM Judges Agree With Each Other and Disagree With Human-Grounded Readability Metrics

Ajay Mandyam Rangarajan*
Independent Researcher
arangarajan@pm.me

Jeyashree Krishnan*
Independent Researcher
jeyashree.krishnan@pm.me

Abstract

Text is increasingly written with language models and increasingly assessed by them, which raises a question. When we ask whether a piece of text is well written, what is the reference, the human reader or the model? We study this for readability, where decades of formulas calibrated on human reading performance supply an external, model independent reference against which a judge can be checked. On a longitudinal corpus of 30,595 NeurIPS abstracts from 1987 to 2025 we compute 15 classical readability metrics, and on the 24,772 of them from 1987 to 2024 we run six open-weight instruction-tuned models under three prompt templates, of which 24,754 returned a parseable score. The 15 classical metrics agree that readability declined over the period. All six models under all three prompts report the opposite, drifting without a consistent direction before 2022 and rising in all 18 model and prompt combinations after it. We then measure the reliability signals a practitioner would actually check. The judges are highly repeatable under resampling, between 44% and 98% of abstracts receive an identical score on all three sampled generations, and they agree with each other at Spearman 0.37 to 0.41 per paper, giving a panel reliability α of 0.75 to 0.79. That is above the 0.70 conventionally taken as acceptable, so on the usual criteria this panel would be reported as reliable. The direction of its yearly trend is also robust to the standardization baseline. Every one of these checks passes, and all six models still point the opposite way from the reference. Agreement and stability are therefore not evidence of validity. This matters wherever a judge is asked to score a subjective human construct, one defined by its effect on a reader, such as readability, cognitive load, clarity or helpfulness, where no ground truth is available to check the judge against. A per-paper analysis shows the judge score correlates positively with density and sensational

language and negatively with sentence length, a partly opposed feature set from the classical reference. If model judgment becomes the operative standard for clear writing, the human reading cost these formulas were built to estimate is no longer measured at any point in the process.

1 Introduction

When we ask whether a piece of text is well written, we are implicitly asking who the reader is. Until recently the answer was uncontroversial, because text was written by people for people, and the standard against which writing was assessed was human comprehension. That premise is weakening on both sides. On the production side, a detector based audit of 43,000 Common Crawl URLs finds machine generated articles at parity with human written ones on the English web (Druck et al., 2025). Measured LLM assistance is climbing in parallel across scientific and institutional text, reaching up to 22% of computer science arXiv abstracts by September 2024 (Liang et al., 2025b). Consumer complaints and corporate press releases stood at roughly 18% and 24% by late 2024 (Liang et al., 2025a). On the assessment side, using a large language model as a judge has become a standard way to score generated and written text, from summary quality to helpfulness (Zheng et al., 2023; Liu et al., 2023; Gu et al., 2026). Text is therefore increasingly written by models and scored by models, and the human reader is progressively less present at either end.

That makes the reliability of the judge a question about more than measurement error. When the judge emits a score, do the usual reliability heuristics, agreement across models, robustness across prompts, and stability across samples, indicate that it is measuring the intended human construct? These heuristics are in effect an uncertainty estimate assembled from the dispersion of a judge panel, and they are routinely read as confidence.

* Equal contribution.




Figure 1: **Overview of the setup.** One corpus of 30,595 NeurIPS abstracts (1987–2025) is scored along two independent paths, 15 classical readability formulas calibrated on human reading performance, and six open-weight instruction-tuned models under three prompt templates, standardized per model. The reference falls (reading gets harder) while the judges are flat and then rise after 2022 (text is judged easier).

We argue that for subjective constructs they can be systematically misleading, and we demonstrate it for readability.

Readability is an unusually good testbed for two reasons. First, it has a mature set of formulas calibrated against human reading performance (Flesch, 1948; Kincaid et al., 1975; DuBay, 2004) and motivated by cognitive load, penalizing long sentences and unfamiliar words. Britton et al. (1982) show simplified syntax reduces the cognitive capacity a reader expends, the mechanism cognitive load theory describes (Sweller, 1994). These provide an external reference that does not depend on any language model. Second, a longitudinal corpus provides a strong reference trend. If a construct moves consistently in one direction over decades under a consensus of validated metrics, a judge that measures the same construct should recover that direction.

We assemble such a corpus, 30,595 NeurIPS abstracts from 1987 to 2025, and compare six open-weight instruction-tuned models under three prompt templates (Cachola et al., 2025) against 15 classical readability metrics (Figure 1). The 15 metrics agree that readability declined over the period. Every model under every prompt reports the opposite. Our contributions are as follows.

- We show a systematic, directional disagreement between LLM readability judges and a 15-metric human-grounded consensus over a 38-year corpus, replicated across six models and three prompts (Section 4).
- We measure the reliability signals themselves, repeatability under resampling, inter-judge

correlation, panel reliability, and sensitivity to the standardization baseline, and show that every one of them passes while the judges are wrong in direction (Sections 4.3 and 4.4).

- We draw the practical implication that reliability for a subjective construct must be validated against an external human-grounded reference, and set out what follows if model judgment displaces the human reader as the standard for clear writing (Section 5).

2 Related work

LLM as judge. Scoring text with a language model is used for summary quality, helpfulness, and open ended generation (Zheng et al., 2023; Liu et al., 2023; Gu et al., 2026). Such judges are reference free, which allows evaluation at scales human annotation does not reach.

Known failure modes. Reported biases are measured at the level of individual judgments. Zheng et al. (2023) identify position bias, verbosity bias, and self enhancement. Panickssery et al. (2024) find that evaluators recognize their own generations and score them higher. Stureborg et al. (2024) report a preference for text of lower perplexity, together with skewed score distributions and sensitivity to immaterial prompt changes. We instead examine whether a panel that passes the aggregate reliability checks in common use recovers a population level trend.

Readability as a measured construct. Classical formulas regress comprehension outcomes on surface features (Flesch, 1948; Kincaid et al., 1975),

with limits documented in the readability literature (DuBay, 2004). Readability of scientific writing has declined over decades on these measures (Plavén-Sigra et al., 2017). Cachola et al. (2025) compare LLM judges with human readability ratings on plain language summaries and report that the judges track those ratings more closely than Flesch-Kincaid does. We adopt their protocol and ask whether a judge recovers the direction of change over time, which their design does not address.

Uncertainty in judge panels. Where a single judge is not trusted, one common practice is to poll several and treat their spread as an uncertainty estimate. We compare that spread with an external reference.

3 Setup

Corpus. We use the abstracts of 30,595 NeurIPS papers from 1987 to 2025, normalized to a single validated schema. Corpus source and the record schema are given with the code and data availability statement. Readability scores are computed on the abstract as published. Markup such as mathematical expressions and URLs is left in place. Re-computing the series with mathematical expressions replaced by a one-syllable placeholder and URLs removed moves the 1987 to 2025 decline from -16.14 to -15.31 Flesch points, so markup accounts for 5.2 percent of it and the reported trend is, if anything, a slight overstatement. The longitudinal span is the design choice that makes the reference trend informative. A single-year snapshot would give no direction to recover.

Reference metrics. We compute 15 classical readability formulas with the `textstat` library. Thirteen are formulas over surface features such as sentence length and word length, several of them fitted to human comprehension outcomes (Flesch, 1948; Kincaid et al., 1975; DuBay, 2004). The remaining two, mean sentence length and mean syllables per word, are the raw counts those formulas take as inputs. We report Flesch Reading Ease (Flesch, 1948) as the primary indicator because it is the only one of the 15 for which higher means easier, so its direction is unambiguous, and we verify that all 15 move together. The 15 are not independent instruments. They share surface inputs, and over NeurIPS their mean absolute pairwise Pearson correlation is 0.68, with the first principal compo-

nent of the standardised metrics accounting for 72% of the variance. Unless stated otherwise, correlations between a judge score and another quantity are Spearman rank correlations, since judge scores are ordinal on a 1 to 5 scale. Appendix B gives every statistic we use. We therefore treat their consensus as one well attested direction of change rather than as 15 confirmations. We do not claim these formulas are good item-level measures. Cachola et al. (2025), whose judge protocol we adopt, report the opposite, that on plain language summaries LLM judges track human ratings better than Flesch-Kincaid. Our claim concerns a different quantity. We use only their consensus on the direction of change over four decades, which item-level accuracy does not entail recovering.

Judges. We follow the LLM as judge readability protocol of Cachola et al. (2025) and apply their three prompt templates, `simple`, `aschb`, and `own_reasoning` (Appendix C), to six open-weight instruction-tuned models: Gemma-3-27B-Instruct, Gemma-4-31B-Instruct, Llama-3.1-8B-Instruct, Mistral-7B-Instruct, Mixtral-8x7B-Instruct, and Qwen2.5-32B-Instruct. Each prompt asks for a 1 to 5 reading-ease rating. Decoding uses temperature 0.7 with three sampled generations per (paper, prompt, model) triple. Each model’s yearly mean is standardized to its own 1987 to 2022 mean and standard deviation, removing absolute scale differences between models, and we report the average of the six z-scored series. These six models were selected to make agreement between them as informative as possible. They span four developers, four model families, both dense and mixture of experts architectures, and roughly 7 to 32 billion active parameters, and they were trained on separately curated corpora. Agreement across a set chosen for that spread is harder to attribute to a single shared pretraining pipeline than agreement within one family. Open weights also allow the protocol to be rerun on the same released models, which scored access to proprietary models does not allow.

4 Results

4.1 Trend in the reference metrics

Figure 2 shows Flesch Reading Ease on NeurIPS from 1987 to 2025. It falls from approximately 24 to approximately 8, a sustained decline across 38 years. All 15 classical metrics move in the same




Figure 2: **Flesch Reading Ease on NeurIPS abstracts, 1987–2025.** Higher values indicate easier reading (↑). NeurIPS falls from approximately 24 in 1987 to approximately 8 in 2025, a sustained decline. All 15 classical metrics move in the same direction on NeurIPS (Appendix A). Vertical dashed line marks the late-2022 availability of widely-used instruction-tuned writing assistants.

direction on NeurIPS. The full 15-metric panels appear in Appendix A.

4.2 Trend in judge scores

Figure 3 shows the standardized judge scores under each of the three prompts. None of the 18 model and prompt combinations reproduces the decline the reference shows. Before 2022 they drift without a consistent direction, and after 2022 every one of them rises. Over the corpus as a whole the judges therefore report readability as unchanged or improving, which contradicts the 15-metric reference in direction.

Figure 4 shows the same data without averaging, one line per model in each prompt panel. For each of the 18 combinations we fit a straight line through its yearly scores and read off the slope, which is the average change in z per year. Before 2022 the slopes point in no consistent direction, nine of the 18 upward and nine downward, and none is steeper than 0.08 z per year, which over three decades is small against a year to year spread of about 4 z . After 2022 all 18 point the same way, rising by between 1.24 and 3.94 z .

4.3 Reliability of the judge panel

We quantify four properties of the judge panel. Comparing models against each other requires abstracts that all six scored, so the agreement statistics use that common set. The other three properties are computed per model and use whatever each model scored.




Figure 3: **LLM as judge readability on NeurIPS abstracts, 1987–2024.** Each curve is the mean z -score across all six open-weight models for one prompt variant, standardized to each model’s own 1987–2022 baseline (dotted line at zero). Under all three prompts the judged readability is approximately flat before 2022 and rises after the availability of instruction-tuned writing assistants (vertical dashed line). The direction is opposite to the classical readability decline in Figure 2.

Repeatability. Decoding is stochastic at temperature 0.7. For each abstract, model, and prompt we compute the standard deviation across the three sampled generations. Its mean ranges from 0.01 to 0.21 on the 1 to 5 scale, and between 44% and 98% of abstracts receive an identical score on all three runs (Table 1).

Agreement between judges. Per-paper Spearman correlation over the 15 model pairs averages 0.37 to 0.41 by prompt, with a minimum of 0.24, and Cronbach’s α over the six judges is 0.75 to 0.79 (Table 2). Agreement on the yearly means is higher than on individual abstracts, since the six series share a common shape (Figure 4).

Resolution of the judge scale. The judges use only a narrow band of the five point scale. Between 97% and 99% of abstracts receive a six-model mean that rounds to 2 or 3, and the middle 98% span 1.33 to 3.25 under `simple`, so in practice the panel resolves roughly two levels rather than five. Two consequences follow. First, the models disagree about level far more than any of them moves over time. Under `ascb` the same abstracts average 1.38 for Gemma-4-31B and 3.55 for Mistral-7B, whereas the 1987 to 2024 change for a single model lies between -0.09 and $+0.44$ points (Table 1). Standardizing each model against its own history removes that level difference and keeps the temporal shape, which is why the series




Figure 4: **Per-model judge scores by prompt template, standardized separately.** Each panel is one prompt template and each line one model, standardized to its own 1987 to 2022 mean and standard deviation. The dashed line marks late 2022. For each combination we fit a straight line through its yearly scores. The slope of that line is the average change in z per year. Before 2022 the slopes point in no consistent direction, nine of the 18 upward and nine downward, and none is steeper than 0.08 z per year, which is small against a year-to-year spread of about 4 z . After 2022 all 18 point the same way, rising by between 1.24 and 3.94 z . The model-averaged series in Figure 3 does not show the pre-2022 level differences or the threefold spread in the size of the rise, which follows from the differing baseline standard deviations.

Model	1987	2024	Δ	SD	Same
Gemma-3-27B	2.14	2.17	+0.03	0.02	98%
Gemma-4-31B	1.57	1.67	+0.10	0.03	95%
Llama-3.1-8B	1.97	2.20	+0.24	0.15	75%
Mistral-7B	2.62	3.04	+0.41	0.21	44%
Mixtral-8x7B	2.81	3.26	+0.44	0.09	83%
Qwen2.5-32B	2.18	2.10	-0.09	0.04	92%

Table 1: **Raw judge scores and run-to-run spread.** Judge score on the 1 to 5 scale under the `simple` prompt, in 1987 and 2024, with the change Δ . SD is the mean standard deviation across the three sampled generations for a given abstract, and Same is the proportion of abstracts scored identically on all three runs. The 1987 column rests on 59 to 87 scored abstracts depending on model, against 4,493 in 2024, so the per-model Δ is estimated far less precisely at the 1987 endpoint than at the 2024 one.

in Figure 3 are large where the raw movement is small. Second, an instrument that resolves two levels has little room to register a slow change, whatever its direction. This is a property of the judge rather than of this corpus, and it is a further reason not to read a flat judge series as evidence that the underlying property held steady.

Sensitivity to the baseline. Re-standardizing on 1987 to 2010, 1987 to 2016, or the full period places the 2024 model-averaged z between 2.48 and 4.47, compared with 3.17 to 3.87 for the 1987

Prompt	mean ρ	min	max	α
<code>simple</code>	0.396	0.286	0.487	0.786
<code>ascb</code>	0.370	0.241	0.478	0.754
<code>own_reasoning</code>	0.409	0.296	0.496	0.791

Table 2: **Inter-judge agreement.** Per-paper Spearman ρ over the 15 model pairs, and Cronbach’s α over the six judges, on the abstracts scored by all six ($n = 23,257$ to 24,609 depending on prompt).

to 2022 baseline. The direction is unchanged under all four windows.

4.4 Agreement and validity

All six models, spanning four families and roughly a factor of five in active parameter count, show the same post-2022 rise. The agreement holds across three prompt formulations, two of which explicitly instruct the model not to use readability formulas and to rely on its own judgment. The judges are repeatable, they correlate with one another, the panel mean is reliable, and the direction survives every baseline we tried. If inter-model agreement, prompt robustness, or sample stability were taken as evidence of reliability, this judge would pass every test. It is nonetheless wrong in direction against a 15-metric human-grounded consensus.

4.5 Association with writing metrics

To characterize the disagreement, we compute the per-paper Spearman correlation between the judge score and every writing metric, over the 24,754 NeurIPS abstracts the judge scored (Figure 5). The judge score covaries with some features of reader effort. It correlates negatively with sentence length and syntactic complexity, as the classical formulas do, but positively with noun phrase density, verb and noun density, and sensational language. Every correlation is small, under 0.25 in absolute value, and the pattern holds across all six models.

5 Discussion

The result bears on LLM-as-a-judge evaluation of subjective human constructs. Dispersion across judges is widely used as the uncertainty estimate attached to a judge score. Here the panel agrees, and it agrees on the wrong direction, so its spread is a poor guide to whether the score is right. Spread measures how much the judges differ from one another. It cannot measure what they get wrong together.

The agreement varies little with model provenance. Models with the spread of developers, families and architectures described in Section 3 would be expected to err at least partly independently. The observed convergence therefore does not follow from one pretraining corpus or architecture. We do not identify the mechanism. What these models do have in common is how they were tuned. Five of the six were trained on ranked human or model preferences (Christiano et al., 2017; Ouyang et al., 2022), increasingly model-ranked as AI feedback replaces human feedback at scale (Bai et al., 2022; Lee et al., 2024), and that objective is known to induce systematic surface preferences that do not correspond to quality, most clearly a preference for length (Singhal et al., 2024). The exception is Mistral-7B-Instruct, which is supervised fine-tuned only, so preference optimization cannot be the whole account. This remains untested here.

The disagreement has measurable structure (Section 4.5), although every correlation is small. The judge co-varies with a different and partly opposed feature set from the classical reference, and every per-paper correlation is under 0.25 in absolute value, against the 15 classical metrics and the remaining metrics alike.

This design measures the direction in which the two instruments disagree. It does not establish

which of them better reflects how readers experience the text, and we do not validate the classical formulas as item-level measures of reading difficulty. On plain language summaries, Cachola et al. (2025) find that their best language model judge correlates with human readability ratings at 0.56 Pearson, ahead of most of the eight traditional metrics they compare. On that basis they recommend reporting model judgments alongside the two best-correlating formulas rather than in place of them. We cannot say which instrument better measures scientific abstracts. Answering that would require readers to rate the same texts. Independent of that determination, six judges agreed closely with one another while moving opposite to a calibrated external measure, so agreement among judges does not establish validity.

The immediate implication for practice is that reliability for a subjective construct must be checked against an external reference rather than against internal consistency. Where such a reference exists, judge scores should be reported alongside it, and any divergence in direction should be reported rather than aggregated away. Where none exists, inter-judge agreement can result from a bias shared across the panel.

An alternative account. The rise after 2022 has an alternative explanation that this design does not separate from judge invalidity. Abstracts written after late 2022 are increasingly LLM assisted (Liang et al., 2025b), and LLM judges are known to favour text they find familiar. Stureborg et al. (2024) document a preference for text the model assigns lower perplexity, and Panickssery et al. (2024) show that evaluators can recognize and favour their own generations. If recent abstracts are partly machine written, a judge with that bias would score them higher without readability having improved on any human measure. This account also predicts the timing of the change.

Either way the measured result stands, because on both accounts the judge score fails to follow the reference. The account also shows what the spread of the panel cannot catch. If all six models share a preference for familiar text, they will agree with one another, and their agreement will look like confidence. Resampling and rewording the prompt do not help either, since both leave the shared preference in place. Separating the two accounts requires abstracts of known provenance, which this corpus does not record.




Figure 5: **Per-paper correlation between writing metrics and judge score.** Spearman correlation between each writing metric and each model’s judge score, over 24,754 NeurIPS abstracts. A model’s score for an abstract is the median of its nine generations, three prompts by three sampled runs. Rows are sorted by the correlation with the mean of those six per-model scores (1987–2024, the years the judge was scored). The judge score runs 1 to 5 with higher meaning easier, so each cell’s sign says which way that feature moves the judge. **Red (positive): more of that feature is judged *more* readable**, as for verb density, sensational language and noun phrase density. **Blue (negative): more of it is judged *less* readable**, as for sentence length and syntactic complexity, which is the direction the classical formulas also take. Near-white cells, including acronym density, indicate near indifference. All values are Spearman correlations and every effect is small ($|r| < 0.25$).

Implications for scientific writing. If text is drafted, filtered and rated by models, as Section 1 indicates it increasingly is, every quality signal in the loop is a model signal, and our result says those signals do not track the human reading cost the classical formulas estimate. Nothing in that process measures the effort a reader spends. The decline we observe began decades before language models were written with or scored by, so we do not attribute it to either. Our claim is narrower. The instruments now being adopted to assess writing quality do not detect that decline, and report it in the opposite direction. Scientific writing is addressed to human readers, so this gap applies directly to it. Evaluations of writing quality should report a human-grounded reference alongside any model score.

Code and data availability

The full pipeline is publicly available at <https://github.com/armandyam/abstract-audit>. The release covers the scraping and schema-validation code for the NeurIPS corpus, together with the per-paper metric scripts for the 15 classical readability formulas and the further writing measures used in the judge comparison. It includes the judge harness, with its three prompt templates and decoding parameters, and the aggregation and plotting scripts that produce every figure in this paper. It also covers the derived per-paper metrics and the 1.3 million judge scores over 24,772 papers.

The abstracts themselves are not redistributed. Each record carries an identifier that resolves to its source page on papers.nips.cc, and the scraping code retrieves the text from there. Abstracts before 2008 are taken from proceedings.neurips.cc, because

the early conference pages bundle the full text into one raw field. Every record is validated against a Pydantic schema carrying paper identifier, venue, year, title, abstract and author list. All six judges are open-weight, and the prompts are reproduced verbatim in Appendix C, so the scoring step can be reproduced independently of our release.

Limitations

The analysis covers abstracts rather than full texts, so the magnitude of the reference trend may differ for full documents, though the direction is unlikely to reverse. [Plavén-Sigra et al. \(2017\)](#) find abstract and full text readability strongly related in a secondary full text analysis. We use six open-weight models with active parameter counts of roughly 7 to 32 billion, and larger closed-weight models may recover more of the reference signal, which we cannot test here. The classical formulas are themselves imperfect proxies for human reading effort. We use only the direction of their consensus and not their absolute values. All 15 formulas could nonetheless share a bias that a language model avoids, and this analysis cannot test that, since no per-paper correlation in Section 4.5 exceeds 0.25. We also cannot separate judge invalidity from the familiarity account discussed in Section 5. Finally, we study a single construct in a single domain, scholarly abstracts. Whether judges of other subjective constructs show the same pattern is untested.

References

- Yuntao Bai, Saurav Kadavath, Sandipan Kundu, Amanda Askell, Jackson Kernion, Andy Jones, Anna Chen, Anna Goldie, Azalia Mirhoseini, Cameron McKinnon, et al. 2022. [Constitutional AI: Harmlessness from AI feedback](#). *Preprint*, arXiv:2212.08073.
- Bruce K. Britton, Shawn M. Glynn, Bonnie J. F. Meyer, and M. J. Penland. 1982. [Effects of text structure on use of cognitive capacity during reading](#). *Journal of Educational Psychology*, 74(1):51–61.
- Isabel Cachola, Daniel Khashabi, and Mark Dredze. 2025. [Evaluating the evaluators: Are readability metrics good measures of readability?](#) In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 24011–24027, Suzhou, China. Association for Computational Linguistics.
- Paul F. Christiano, Jan Leike, Tom B. Brown, Miljan Martic, Shane Legg, and Dario Amodei. 2017. [Deep reinforcement learning from human preferences](#). In *Advances in Neural Information Processing Systems* 30, pages 4299–4307.
- Gregory Druck, Jose Luis Paredes, Bevin Benson, and Ethan Smith. 2025. [More articles are now created by AI than humans](#). Graphite industry report. Detector based estimate over 43,000 Common Crawl URLs sampled from January 2020 to May 2025.
- William H. DuBay. 2004. *The Principles of Readability*. Impact Information, Costa Mesa, CA.
- Rudolph Flesch. 1948. [A new readability yardstick](#). *Journal of Applied Psychology*, 32(3):221–233.
- Jiawei Gu, Xuhui Jiang, Zhichao Shi, Hexiang Tan, Xuehao Zhai, Chengjin Xu, Wei Li, Yinghan Shen, Shengjie Ma, Honghao Liu, Saizhuo Wang, Kun Zhang, Zhouchi Lin, Bowen Zhang, Lionel Ni, Wen Gao, Yuanzhuo Wang, and Jian Guo. 2026. [A survey on LLM-as-a-judge](#). *The Innovation*, 7(6):101253.
- J. Peter Kincaid, Robert P. Fishburne, Richard L. Rogers, and Brad S. Chissom. 1975. Derivation of new readability formulas (automated readability index, fog count and flesch reading ease formula) for Navy enlisted personnel. Technical Report Research Branch Report 8–75, Naval Air Station Memphis, Research Branch.
- Harrison Lee, Samrat Phatale, Hassan Mansoor, Thomas Mesnard, Johan Ferret, Kellie Lu, Colton Bishop, Ethan Hall, Victor Carbune, Abhinav Rastogi, and Sushant Prakash. 2024. [RLAIF vs. RLHF: Scaling reinforcement learning from human feedback with AI feedback](#). In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 26874–26901. PMLR.
- Weixin Liang, Yaohui Zhang, Mihai Codreanu, Jiayu Wang, Hancheng Cao, and James Zou. 2025a. [The widespread adoption of large language model-assisted writing across society](#). *Patterns*, 6(12):101366.
- Weixin Liang, Yaohui Zhang, Zhengxuan Wu, Haley Lepp, Wenlong Ji, Xuandong Zhao, Hancheng Cao, Sheng Liu, Siyu He, Zhi Huang, Diyi Yang, Christopher Potts, Christopher D. Manning, and James Zou. 2025b. [Quantifying large language model usage in scientific papers](#). *Nature Human Behaviour*, 9(12):2599–2609.
- Yang Liu, Dan Iter, Yichong Xu, Shuohang Wang, Ruochen Xu, and Chenguang Zhu. 2023. [G-Eval: NLG evaluation using GPT-4 with better human alignment](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 2511–2522.
- Long Ouyang, Jeff Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul F. Christiano, Jan Leike, and Ryan Lowe. 2022. [Training language models to follow instructions with human feedback](#). In *Advances in Neural Information Processing Systems* 35, pages 27730–27744.

Arjun Panickssery, Samuel R. Bowman, and Shi Feng. 2024. [LLM evaluators recognize and favor their own generations](#). In *Advances in Neural Information Processing Systems* 37.

Pontus Plavén-Sigraý, Granville James Matheson, Björn Christian Schiffler, and William Hedley Thompson. 2017. [The readability of scientific texts is decreasing over time](#). *eLife*, 6:e27725.

Prasann Singhal, Tanya Goyal, Jiacheng Xu, and Greg Durrett. 2024. [A long way to go: Investigating length correlations in RLHF](#). In *First Conference on Language Modeling (COLM)*.

Rickard Stureborg, Dimitris Alikaniotis, and Yoshi Suhara. 2024. [Large language models are inconsistent and biased evaluators](#). *Preprint*, arXiv:2405.01724.

John Sweller. 1994. [Cognitive load theory, learning difficulty, and instructional design](#). *Learning and Instruction*, 4(4):295–312.

Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhonghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. 2023. [Judging LLM-as-a-judge with MT-bench and Chatbot Arena](#). In *Advances in Neural Information Processing Systems 36: Datasets and Benchmarks Track*.

A Additional figures

Figures 6 and 7 show all 15 classical readability metrics on NeurIPS. Every metric moves in the harder-to-read direction over the period. Each panel is one metric, with the year on the horizontal axis. An upward arrow in a panel title marks the one metric, Flesch Reading Ease, for which higher means easier. Every other panel is a grade level or a difficulty count, for which higher means harder. The two figures split the 15 metrics only so that each panel stays legible.

B Statistics reported

Reported means are arithmetic. Reported standard deviations are sample standard deviations, with denominator $n - 1$.

Collinearity of the 15 reference metrics is summarised in two ways. The first is the mean absolute Pearson correlation across all 105 metric pairs. The second is the share of variance carried by the first principal component of the standardised metrics. Both are computed on the papers that have a complete set of scores.

Judge agreement. Judge scores are ordinal on a 1 to 5 scale, so agreement between two judges a and b is the Spearman rank correlation ρ_{ab} , computed over the papers both scored. Panel agreement is the mean of ρ_{ab} over all distinct judge pairs within a prompt.

Standardisation of judge scores. To place six judges on one scale, each model’s annual mean is expressed in units of that model’s own variation over its pre-2022 baseline years at the same venue,

$$z_{m,t} = \frac{\mu_{m,t} - \mu_{m,\text{base}}}{s_{m,\text{base}}},$$

where $\mu_{m,t}$ is model m ’s mean score in year t , and $\mu_{m,\text{base}}$ and $s_{m,\text{base}}$ are the mean and standard deviation of $\mu_{m,\cdot}$ over the baseline years, 2022 included. The reported curve is the unweighted mean of $z_{m,t}$ over the six models.

Association with writing metrics. Per-paper association between a judge score and a writing metric is again Spearman. The heatmap reports it pooled across years. The two summary columns give both the pooled value and a within-year value, computed on scores standardised inside each venue-year cohort so that the trend common to both series cannot induce the correlation. The two agree to within 0.03 in absolute value.

C Prompt templates

Reproduced verbatim from [Cachola et al. \(2025\)](#), Table 10, except that the Reason: <reasoning> output line is omitted because we take only the numeric score. {SUMMARY} is replaced by the abstract at inference.

simple

On a scale of 1 to 5, what is the reading ease of the following text? 1 indicates the text requires expert background knowledge and 5 indicates the text is readable to the general population. Assume the reader is an adult.

Format the output as follows:

Score: <score>

Text: {SUMMARY}




Figure 6: **Classical readability metrics on NeurIPS, part A: Flesch family and grade-level formulas, 1987–2025.** Arrows indicate the direction of easier reading ($\uparrow$ = higher is better, $\downarrow$ = lower is better). Every grade-level metric rises on NeurIPS after 2015 and accelerates after 2020, and Flesch Reading Ease drops correspondingly. Vertical dashed lines mark the late-2022 availability of instruction-tuned writing assistants. Part B with the remaining seven metrics appears in Figure 7.

ascb On a scale of 1 to 5, what is the reading ease of the following text? 1 indicates the text requires expert background knowledge and 5 indicates the text is readable to the general population. Characteristics of a highly readable text include: - Know your audience, and focus and organize your information for that particular audience. - Focus on the big picture. What larger problem is your work a part of? What major ideas or issues does your work address? How will your work help global understanding of some issue? - Avoid jargon. If you must use a technical term, make sure to explain it, but simplify the language. - Try to use metaphors or analogies to	everyday experiences that people can relate to. - Underscore the importance of public support for exploratory research and scientific information, and the role of this information in providing the context for effective policy making. Assume the reader is an adult. Do not use Flesch-Kincaid or other readability formulas. Use your own judgment to rate the text. Format the output as follows: Score: <score> Text: {SUMMARY}
	own_reasoning On a scale of 1 to 5, what is the reading ease of the following text? 1 indicates the text requires expert




Figure 7: **Classical readability metrics on NeurIPS, part B: Linsear Write, LIX, RIX, FORCAST, Powers-Sumner-Kearl, sentence length, and syllables per word, 1987–2025.** Arrows indicate the direction of easier reading ($\downarrow$ = lower is better). Vertical dashed lines mark the late-2022 availability of instruction-tuned writing assistants. Part A with the Flesch family and the first eight grade-level formulas appears in Figure 6.

background knowledge and 5 indicates the text is readable to the general population.

Assume the reader is an adult. Do not use Flesch-Kincaid or other readability formulas. Use your own judgment to rate the text.

Format the output as follows:

Score: <score>

Text: {SUMMARY}